


# Fuzzing and Browser Exploits

MST - 2/23



# Memory Corruption

- **Memory Corruption** attacks occur when we can write somewhere to memory we aren't supposed to
- We can exploit **memory structures** to change the flow of the program
- Methods we've discussed
  - Shellcoding
  - ret2libc
  - ROP
- What else is there?
  - Lots!
  - ... but we won't get into them :(

# Other (common) forms of Memory Corruption

- Heap overflows
- Integer overflows
- Format string attacks
- Use-after-free
- Double free
- Double fetch
- SROP

# Integer Overflows

- Huge or negative numbers are passed in and improperly handled
- Often used to bypass checks or to write memory past the bounds of a buffer
- What's wrong with this program?

```
#include <stdlib.h>

int main(int argc, char *argv[]) {
    int num = atoi(argv[1]);
    char *input = malloc(4 * num);
    for(int x = 0; x < num; x++) {
        input[4*x] = x;
        input[4*x+1] = x;
        input[4*x+2] = x;
        input[4*x+3] = x;
    }
    /* do some other stuff */
    return 0;
}
```

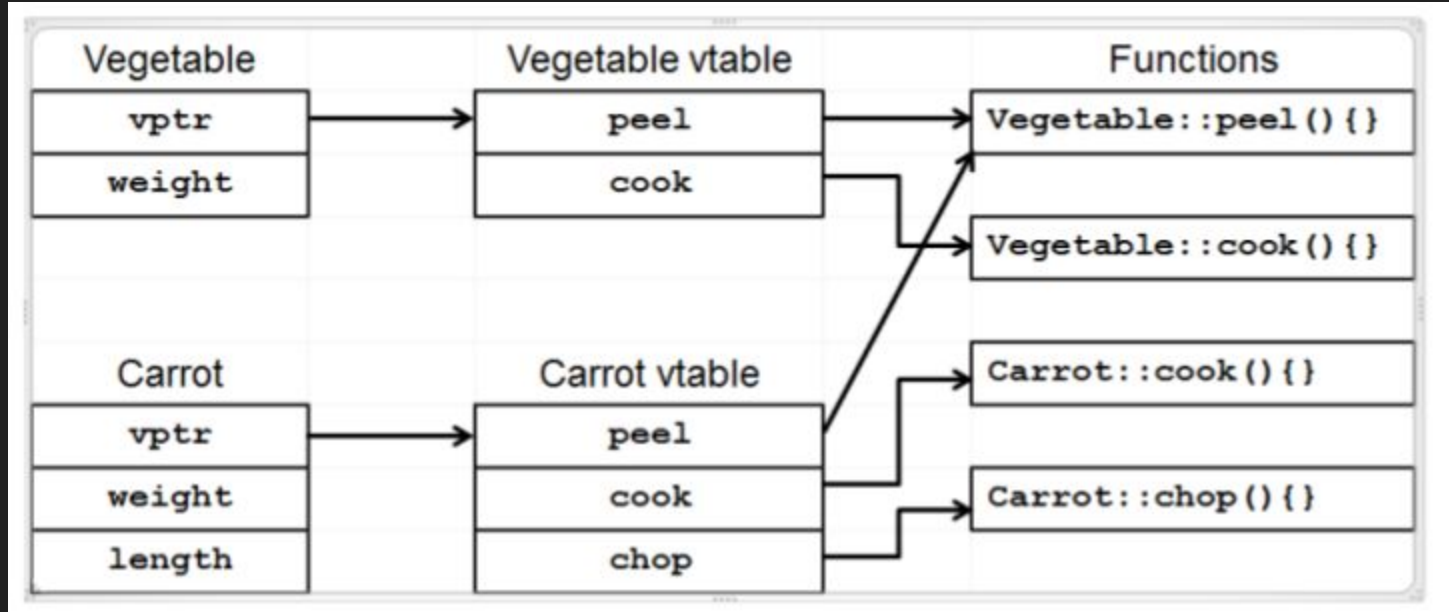
# Heap overflows

- Overflowing a data structure, but on the heap instead of the stack
- Heap is for dynamically allocated memory
  - eg. `malloc(100)` goes on the heap
- We can overwrite from one heap object to another nearby one and change the contents of data

|  |           |   |   |   |   |  |          |  |  |  |        |   |   |   |   |
|--|-----------|---|---|---|---|--|----------|--|--|--|--------|---|---|---|---|
|  | Our Array |   |   |   |   |  | (unused) |  |  |  | Target |   |   |   |   |
|  | C         | Y | R | U | S |  |          |  |  |  | H      | E | L | L | O |

[illegible]

# Use-after-free



NCC Group, *Exploiting CVE-2014-0282*

# Double Free

- A pointer's contents may change after it has been freed
- Calling free on corrupted heap metadata can cause your program to crash
- Calling free on attacker-controlled heap metadata can cause your program to be exploited
- Set your freed pointers to null

```
#include <stdio.h>
#include <unistd.h>
#define BUFSIZE1 512
#define BUFSIZE2 ((BUFSIZE1/2) - 8)

int main(int argc, char **argv) {
    char *buf1;
    char *buf2;
    char *buf3;
    buf1 = (char *) malloc(BUFSIZE2);
    buf2 = (char *) malloc(BUFSIZE2);
    free(buf1);
    free(buf2);
    buf3 = (char *) malloc(BUFSIZE1);
    strncpy(buf3, argv[1], BUFSIZE1-1);
    free(buf2);
    free(buf3);
}
```



# Fuzzing

- Security and binary defenses are improving
  - Time + money investment to find a single bug is increasing
- Exploitation is (usually) straightforward -- finding the bug is hard
- How can we efficiently test large software with thousands of functions?
  - It takes long enough for even small programs...
- **Fuzzing** lets us test software automatically for vulnerabilities
  - Responsible for **overwhelming majority** of all new software vulnerabilities found in major products

# Fuzzing

1. Take a series of valid program inputs
2. For each input string in the list, perform **random mutations** on it to generate a new input
3. Send the new input to the program and see if it crashes
4. Repeat forever!

[illegible]

# Fuzzing

1. Take a series of valid program inputs
2. For each input string in the list, perform **random mutations** on it to generate a new input
3. Send the new input to the program and see if it crashes
4. Repeat forever!

```
void do_things(int which) {  
    void (*call_me)();  
    switch(which) {  
        case 1:  
            call_me = foo;  
            break;  
        case 2:  
            call_me = bar;  
            break;  
    }  
    call_me();  
}
```

# “Dumb” Fuzzers

- No concept of data format, only as an input string
- No setup or prior-knowledge of the target involved
- Randomly mutations some known inputs
  - Flip bits
  - Flip bytes
  - Switch bytes
  - Add or remove random data

```
#include <stdio.h>
int main(int argc, char *argv[]) {
    char buffer[10] = {0};
    read(STDIN, buffer, 9);
    if (buffer[3] != 'q') {
        crash_and_burn();
    }
    print("Ok!\n");
    return 0;
}
```

```
$ cat "AAAqAAAAA" | ./weird
Ok!
$ cat "AAAAAAAAA" | ./weird
crash and burn?
```

# “Dumb” Fuzzers

- Very slow for **structured data**  
(Which is most data)
  - File formats, network protocols, data with checksums, etc
- We spent most of the time stuck on the highlighted section, so we may never find the crash!

```
#include <stdio.h>
int main(int argc, char *argv[]) {
    char buffer[10] = {0};
    read(STDIN, buffer, 9);
    if (sum(buffer, 10) != 334) {
        return 0;
    }
    if (buffer[3] != 'q') {
        crash_and_burn();
    }
    print("Ok!\n");
    return 0;
}
```

# Protocol Fuzzers

- Randomly mutate data within the given data structure
- Given a data format, generate **conformant** random input for the program
- Requires **prior knowledge** of the specific file format to work
- Can better target file parsers, network servers/clients, etc

```
#include <stdio.h>
int main(int argc, char *argv[]) {
    char buffer[10] = {0};
    read(STDIN, buffer, 9);
    if (sum(buffer, 10) != 334) {
        return 0;
    }
    if (buffer[3] != 'q') {
        crash_and_burn();
    }
    print("Ok!\n");
    return 0;
}
```

# Protocol Fuzzers

- We still have serious problems
  - We can miss bugs that are valid but **extremely uncommon**
  - We don't know **how much** of the program we have targeted
  - We need a **robust and complete format definition** to even get started



# Evolutionary Fuzzing (AFL)

- Add instrumentation to the program to **track code coverage**
  - Track program branches
- **Keep** input that increases our code coverage and randomly mutate that
- **Discard** input that doesn't improve coverage

```
#include <stdio.h>
int main(int argc, char *argv[]) {
    char buffer[10] = {0};
    read(STDIN, buffer, 9);
    if (buffer[0] == 'a') {
        if (buffer[1] == 'b') {
            if (buffer[2] == 'c') {
                die_hard();
            }
        }
    }
    print("Ok!\n");
    return 0;
}
```



# Evolutionary Fuzzing (AFL)

|                                       |                                      |                                   |
|---------------------------------------|--------------------------------------|-----------------------------------|
| IJG jpeg <sup>1</sup>                 | libjpeg-turbo <sup>1 2</sup>         | libpng <sup>1</sup>               |
| libtiff <sup>1 2 3 4 5</sup>          | mozjpeg <sup>1</sup>                 | PHP <sup>1 2 3 4 5</sup>          |
| Mozilla Firefox <sup>1 2 3 4</sup>    | Internet Explorer <sup>1 2 3 4</sup> | Apple Safari <sup>1</sup>         |
| Adobe Flash / PCRE <sup>1 2 3 4</sup> | sqlite <sup>1 2 3 4...</sup>         | OpenSSL <sup>1 2 3 4 5 6 7</sup>  |
| LibreOffice <sup>1 2 3 4</sup>        | poppler <sup>1</sup>                 | freetype <sup>1 2</sup>           |
| GnuTLS <sup>1</sup>                   | GnuPG <sup>1 2 3 4</sup>             | OpenSSH <sup>1 2 3</sup>          |
| PuTTY <sup>1 2</sup>                  | ntpd <sup>1</sup>                    | nginx <sup>1 2 3</sup>            |
| bash (post-Shellshock) <sup>1 2</sup> | tcpdump <sup>1 2 3 4 5 6 7 8 9</sup> | JavaScriptCore <sup>1 2 3 4</sup> |

# Problems with fuzzing

- Extremely slow for large programs
- We can't know if we've completely tested all possible code
  - We can *try* to reach all code, but that is not the same
- We struggle with checks like `(x == y && y == 0x12345)` that require specific input values to bypass
  - If x and y are 4-byte ints, AFL will try all  $(2^{32})^2$  attempts
  - This is known as the **state explosion** problem
- We gain no information about program meaning

# The Future: SAGE

- Microsoft Research project to incorporate symbolic evaluation and fuzzing
  - “Scalable, Automated, Guided Execution”
- First true “smart” fuzzing
- Analyze the program to find **constraints** on input that drive the fuzzing engine
  - Significant reduction in search space
- Within those constraints, fuzz away!
- Code coverage as measure of progress
- Responsible for  $\frac{1}{3}$  of **pre-release Windows 7 bugs**
- <https://www.microsoft.com/en-us/springfield/>

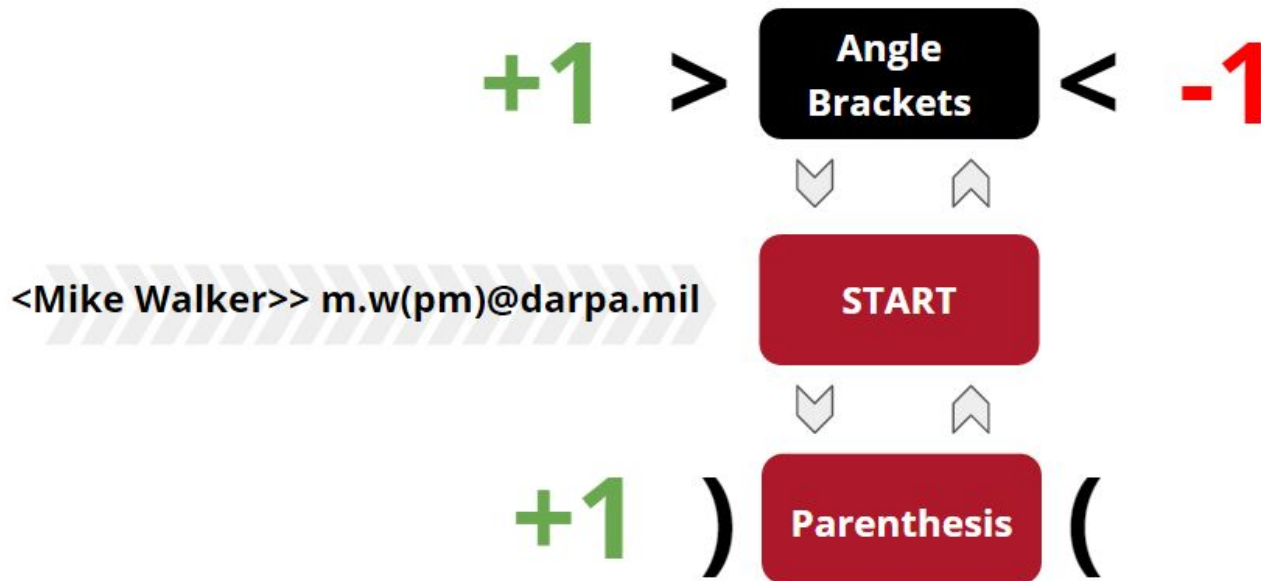
```
#include <stdio.h>
int main(int argc, char *argv[]) {
    int buf[3] = {0};
    read(STDIN, buf, 12);
    if (buf[0] + buf[1] = 993) {
        if (buf[1] - 54 = buf[3]) {
            if (sum(buf) = 82) {
                die_hard();
            }
        }
    }
    print("Ok!\n");
    return 0;
}
```

# The Future: The Cyber Grand Challenge

- DARPA competition to advance the field of fuzzing and program analysis
- 7 finalists developed systems that combined fuzzing, symbolic execution, static analysis, etc
- Led to some advances in the state of the art
- UVa team placed 2nd!



# crackaddr



- **201 iterations** through a loop before the bug has an impact
- $10^{65535}$  possible function inputs, only 1 is vulnerable
- **CGC software found the bug!**

# Hack Firefox

# Adobe Flash CVE-2014-0556

```
public function copyPixelsToByteArray(rect:Rectangle, data:ByteArray) {  
    if(data.position + rect.width * rect.height >= data.length) {  
        // catch buffer overflow  
        return;  
    }  
    // copy data starting at data.position  
    ...  
}
```

# Adobe Flash CVE-2014-0556

- Integer overflow when `data.position` is almost  $2^{32}$
- The out-of-bounds write will be written before the start of the buffer because the pointer will wrap

```
bitmap = new BitmapData(1024, 1024);  
rect = new Rectangle(0, 0, 1024, 1024);  
array = new ByteArray();  
array1.length = 0x40000000;  
array1.position = 0xffffffff000;  
bitmap.copyPixelsToByteArray(rect, array1);
```

- Writeup:  
<https://googleprojectzero.blogspot.com/2014/09/exploiting-cve-2014-0556-in-flash.html>



# A wild 0-day appears

- Someone posts an 0day to the TOR mailing list  
<https://lists.torproject.org/pipermail/tor-talk/2016-November/042639.html>
- Exploit source code analysis  
<https://community.rapid7.com/community/metasploit/blog/2016/12/29/a-friendly-fireside-foray-into-a-firefox-fracas>

# Privileged Javascript Injection

```
1 | crmfObject = crypto.generateCRMFRequest("requestedDN", "regToken", "authenticator",  
2 |                                           "escrowAuthorityCert", "CRMF Generation Done Code",  
3 |                                           keySize1, "keyParams1", "keyGenAlg1",  
4 |                                           ...,  
5 |                                           keySizeN, "keyParamsN", "keyGenAlgN");
```

| Argument                           | Description   |
|------------------------------------|---|
| <i>"CRMF Generation Done Code"</i> | This parameter is JavaScript to execute when the CRMF generation is complete. |

Let's force a call to generateCRMFRequest from a privileged context

# Privileged Javascript Injection

```
1  var y = {};  
2  y.constructor.prototype.toString = function() {  
3      crypto.generateCRMFRequest(..., run_payload, 1024, null, "rsa-ex");  
4      return 5;  
5  };  
6  console.time(y);
```

`console.time` runs in a privileged browser context

# CVE-2013-1710

```
1.1 --- a/security/manager/ssl/src/nsCrypto.cpp
1.2 +++ b/security/manager/ssl/src/nsCrypto.cpp
1.3 @@ -1916,17 +1916,17 @@ nsCrypto::GenerateCRMFRequest(nsIDOMCRMF
1.4     return NS_ERROR_FAILURE;
1.5 }
1.6 jsString = JS_ValueToString(cx, argv[4]);
1.7 NS_ENSURE_TRUE(jsString, NS_ERROR_OUT_OF_MEMORY);
1.8 argv[4] = STRING_TO_JSVAL(jsString);
1.9 JSAutoByteString jsCallback(cx, jsString);
1.10 NS_ENSURE_TRUE(!jsCallback, NS_ERROR_OUT_OF_MEMORY);
1.11
1.12 - nrv = xpc->WrapNative(cx, ::JS_GetGlobalObject(cx),
1.13 + nrv = xpc->WrapNative(cx, JS_GetGlobalForScopeChain(cx),
1.14     static_cast<nsIDOMCrypto *>(this),
1.15     NS_GET_IID(nsIDOMCrypto), getter_AddRefs(holder));
1.16 NS_ENSURE_SUCCESS(nrv, nrv);
1.17
```

JS\_GetGlobalForScopeChain() returns the global object for *whatever function is currently running* on the context.

# CVE-2013-1670

```
2.12 +template <class Base>
2.13 +bool
2.14 +SecurityWrapper<Base>::defineProperty(JSContext *cx, HandleObject wrapper,
2.15 +                                     HandleId id, PropertyDescriptor *desc)
2.16 +{
2.17 +   if (desc->getter || desc->setter) {
2.18 +       JSString *str = IdToString(cx, id);
2.19 +       const jschar *prop = str ? str->getCharsZ(cx) : NULL;
2.20 +       JS_ReportErrorNumberUC(cx, js_GetErrorMessage, NULL,
2.21 +                             JSMMSG_ACCESSOR_DEF_DENIED, prop);
2.22 +       return false;
2.23 +   }
2.24 +
2.25 +   return Base::defineProperty(cx, wrapper, id, desc);
2.26 +}
```

allows remote attackers to bypass certain read-only restrictions

# Homework

# Homework

- Fuzz a simple binary: fuzz\_me.c

## Tips

- **Don't hesitate to ask questions! (Slack, email, etc)**

# Homework grading

- Submit fuzzing results to [cm7bv@virginia.edu](mailto:cm7bv@virginia.edu) with the subject “MST Assignment 5 - <YOUR\_UVA\_ID>”
  - eg: “MST Assignment 5 - cm7bv”
- Also, include a brief (1-paragraph) description of what you did and how it went



# Useful Resources

- The Smart Fuzzer Revolution  
(<https://www.youtube.com/watch?v=g1E2Ce5cBhI>)
- More about crackaddr  
(<http://2015.hackitoergosum.org/slides/HES2015-10-29%20Cracking%20Sendmail%20crackaddr.pdf>)
- DARPA video about UVA's CGC team  
(<https://www.youtube.com/watch?v=CHdmYY-kyuA>)
- More about heap exploits  
(<http://www.mathyvanhoef.com/2013/02/understanding-heap-exploiting-heap.html>)